

Final Report

Nicholas Allen, Surya Maddali, Jake Adams

Introduction:

In recent decades, cell phones have become a hot commodity around the world. The idea of calling with the tips of your fingers was a revolutionary idea that continues to set the standard for telecommunications. With advancements to cell phones, one question that is always present is pricing. There may be certain factors that affect cell phone pricing such as storage, camera capabilities, and battery power. The goal of this project is to assess that, seeing if certain features of phones affect pricing in a significant way. It is an interesting and important question to answer because it can inform others about what phone features matter the most to companies that make phones as well as inform us about what features matter the most to a phone's functionality when looking to buy one. Machine learning is a reasonable approach to tackle this question because it can give us insight into why or how phones are purchased. Moreover, it can help us predict phone prices in the future based on what features they possess, which would be informed by past data on this exact matter. In other words, it would help readers assess what features are continuing to affect the price of the phone the most in the present.

Illustration:

Background and Related Works:

We looked at an article from IEEE Xplore. This article was about predicting mobile phone prices using a data set from kaggle. This article differed from ours because they were predicting phone prices with classification. They had their y variable in as a factor with 4 levels. The levels were from "low cost" to "very high cost". Some examples of their x variables were battery power and clock speed. They used several different models to predict phone price such as a decision tree and SVM. Their most accurate model was SVM with an accuracy of 94.8%.




Figure 1: Elephant

Reference: N. Hu, “Classification of Mobile Phone Price Dataset Using Machine Learning Algorithms,” 2022 3rd International Conference on Pattern Recognition and Machine Learning (PRML), Chengdu, China, 2022, pp. 438-443, doi: 10.1109/PRML56267.2022.9882236.
 keywords: {Support vector machines;Machine learning algorithms;Random access memory;Machine learning;Feature extraction;Mobile handsets;Batteries;computer science;machine learning;classification;price prediction},

Data Processing:

We loaded in the data sets though the readxl package

Loading required package: Matrix

Loaded glmnet 4.1-8

```

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.1      v tibble     3.2.1
  
```

```

v lubridate 1.9.3      v tidyr      1.3.1
v purrr      1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x tidyr::expand() masks Matrix::expand()
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
x tidyr::pack()    masks Matrix::pack()
x tidyr::unpack() masks Matrix::unpack()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become explicit
corrplot 0.92 loaded

```

Warning: package 'caret' was built under R version 4.3.2

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

```

Rows: 3114 Columns: 12-- Column specification -----
Delimiter: ","
chr (7): Brands, Models, Colors, Memory, Storage, Camera, Mobile
dbl (5): Rating, Selling Price, Original Price, Discount, discount percentage
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

First Dataset

The first data set looked like this before processing.

```

# A tibble: 6 x 11
  model      price rating sim  processor ram  battery display camera card  os
  <chr>    <chr>  <dbl> <chr> <chr>      <chr> <chr>   <chr>   <chr> <chr> <chr>
1 OnePlus~  54,~      89 Dual~ Snapdrag~ 12 G~ 5000 m~ 6.7 in~ 50 MP~ Memo~ Andr~
2 OnePlus~  19,~      81 Dual~ Snapdrag~ 6 GB~ 5000 m~ 6.59 i~ 64 MP~ Memo~ Andr~
3 Samsung~  16,~      75 Dual~ Exynos 1~ 4 GB~ 5000 m~ 6.6 in~ 50 MP~ Memo~ Andr~
4 Motorol~  14,~      81 Dual~ Snapdrag~ 6 GB~ 5000 m~ 6.55 i~ 50 MP~ Memo~ Andr~
5 Realme ~  24,~      82 Dual~ Dimensit~ 6 GB~ 5000 m~ 6.7 in~ 108 M~ Memo~ Andr~
6 Samsung~  16,~      80 Dual~ Snapdrag~ 6 GB~ 5000 m~ 6.6 in~ 50 MP~ Memo~ Andr~

```

It is a tabular data set on some mobile phones. Some examples of columns in the data set are mobile which represents the name of the phone and the price of the phone.

To start off we took out the model column because it represented the names of the phones which will not impact the price. We also took out the sim column.

```
# A tibble: 6 x 9
  price rating processor      ram battery display camera card os
  <chr>   <dbl> <chr>          <chr> <chr>   <chr>   <chr> <chr> <chr>
1  54,999    89 Snapdragon 8 Gen2, Oc~ 12 G~ 5000 m~ 6.7 in~ 50 MP~ Memo~ Andr~
2  19,989    81 Snapdragon 695, Octa ~ 6 GB~ 5000 m~ 6.59 i~ 64 MP~ Memo~ Andr~
3  16,499    75 Exynos 1330, Octa Cor~ 4 GB~ 5000 m~ 6.6 in~ 50 MP~ Memo~ Andr~
4  14,999    81 Snapdragon 695, Octa~ 6 GB~ 5000 m~ 6.55 i~ 50 MP~ Memo~ Andr~
5  24,999    82 Dimensity 1080, Octa ~ 6 GB~ 5000 m~ 6.7 in~ 108 M~ Memo~ Andr~
6  16,999    80 Snapdragon 750G, Oct~ 6 GB~ 5000 m~ 6.6 in~ 50 MP~ Memo~ Andr~
```

Cleaning battery column

We extracted the battery life of each phone in mAH and made the column numeric

```
Warning: There was 1 warning in `mutate()`.
i In argument: `battery = .Primitive("as.double")(battery)`.
Caused by warning:
! NAs introduced by coercion
```

```
# A tibble: 6 x 9
  price rating processor      ram battery_mAh display camera card os
  <chr>   <dbl> <chr>          <chr>   <dbl> <chr>   <chr> <chr> <chr>
1  54,999    89 Snapdragon 8 Gen2~ 12 G~      5000 6.7 in~ 50 MP~ Memo~ Andr~
2  19,989    81 Snapdragon 695, 0~ 6 GB~      5000 6.59 i~ 64 MP~ Memo~ Andr~
3  16,499    75 Exynos 1330, Octa~ 4 GB~      5000 6.6 in~ 50 MP~ Memo~ Andr~
4  14,999    81 Snapdragon 695, ~ 6 GB~      5000 6.55 i~ 50 MP~ Memo~ Andr~
5  24,999    82 Dimensity 1080, 0~ 6 GB~      5000 6.7 in~ 108 M~ Memo~ Andr~
6  16,999    80 Snapdragon 750G,~ 6 GB~      5000 6.6 in~ 50 MP~ Memo~ Andr~
```

Cleaning processor variable

We extracted the power of the processor in GHz. We then made the column numeric

```
# A tibble: 6 x 9
```

	price	rating	`processor GHz)`	ram	battery_mAh	display	camera	card	os
	<chr>	<dbl>		<dbl> <chr>	<dbl>	<chr>	<chr>	<chr>	<chr>
1	54,999	89		3.2 12 GB ~	5000	6.7 in~	50 MP~	Memo~	Andr~
2	19,989	81		2.2 6 GB R~	5000	6.59 i~	64 MP~	Memo~	Andr~
3	16,499	75		2.4 4 GB R~	5000	6.6 in~	50 MP~	Memo~	Andr~
4	14,999	81		2.2 6 GB R~	5000	6.55 i~	50 MP~	Memo~	Andr~
5	24,999	82		2.6 6 GB R~	5000	6.7 in~	108 M~	Memo~	Andr~
6	16,999	80		2.2 6 GB R~	5000	6.6 in~	50 MP~	Memo~	Andr~

Cleaning os column

We noticed that because the data was unclean, some of the values that should be in the os column were in the card column. We put these value in the os column and removed the card column after. We also made the os column a factor.

```
# A tibble: 6 x 8
```

	price	rating	`processor GHz)`	ram	battery_mAh	display	camera	os
	<chr>	<dbl>		<dbl> <chr>	<dbl>	<chr>	<chr>	<fct>
1	54,999	89		3.2 12 GB RAM, 2~	5000	6.7 in~	50 MP~	Andr~
2	19,989	81		2.2 6 GB RAM, 12~	5000	6.59 i~	64 MP~	Andr~
3	16,499	75		2.4 4 GB RAM, 64~	5000	6.6 in~	50 MP~	Andr~
4	14,999	81		2.2 6 GB RAM, 12~	5000	6.55 i~	50 MP~	Andr~
5	24,999	82		2.6 6 GB RAM, 12~	5000	6.7 in~	108 M~	Andr~
6	16,999	80		2.2 6 GB RAM, 12~	5000	6.6 in~	50 MP~	Andr~

Cleaning camera column

We extracted the amount of mega pixels in the front camera of each phone. We made this column numeric

```
# A tibble: 6 x 8
```

	price	rating	`processor GHz)`	ram	battery_mAh	display	f_camera_MP	os
	<chr>	<dbl>		<dbl> <chr>	<dbl>	<chr>	<dbl>	<fct>
1	54,999	89		3.2 12 GB R~	5000	6.7 in~	16	Andr~
2	19,989	81		2.2 6 GB RA~	5000	6.59 i~	16	Andr~
3	16,499	75		2.4 4 GB RA~	5000	6.6 in~	13	Andr~
4	14,999	81		2.2 6 GB RA~	5000	6.55 i~	16	Andr~
5	24,999	82		2.6 6 GB RA~	5000	6.7 in~	16	Andr~
6	16,999	80		2.2 6 GB RA~	5000	6.6 in~	8	Andr~

Cleaning ram column

We extracted the ram of the phones in GB and made it a factor because phones only have a few preset values for their ram

```
# A tibble: 6 x 8
  price rating `processor GHz)` ram battery_mAh display f_camera_MP os
  <chr>   <dbl>          <dbl> <fct>      <dbl> <chr>          <dbl> <fct>
1 54,999    89          3.2 12 GB      5000 6.7 inch~      16 Andr~
2 19,989    81          2.2 6 GB       5000 6.59 inch~     16 Andr~
3 16,499    75          2.4 4 GB       5000 6.6 inch~     13 Andr~
4 14,999    81          2.2 6 GB       5000 6.55 inch~     16 Andr~
5 24,999    82          2.6 6 GB       5000 6.7 inch~     16 Andr~
6 16,999    80          2.2 6 GB       5000 6.6 inch~      8 Andr~
```

Cleaning Display column

We extracted the display size and the Hz of the display and turned that into two new columns. We made these new columns numeric and removed the original

```
# A tibble: 6 x 9
  price rating `processor GHz)` ram battery_mAh f_camera_MP os displaySize
  <chr>   <dbl>          <dbl> <fct>      <dbl>      <dbl> <fct>      <dbl>
1 54,9~    89          3.2 12 GB      5000        16 Andr~      6.7
2 19,9~    81          2.2 6 GB       5000        16 Andr~      6.59
3 16,4~    75          2.4 4 GB       5000        13 Andr~      6.6
4 14,9~    81          2.2 6 GB       5000        16 Andr~      6.55
5 24,9~    82          2.6 6 GB       5000        16 Andr~      6.7
6 16,9~    80          2.2 6 GB       5000         8 Andr~      6.6
# i 1 more variable: displayHz <dbl>
```

Cleaning Price column

We converted the value in rupees to dollars to make it easier to understand for our audience

```
# A tibble: 6 x 9
  price rating `processor GHz)` ram battery_mAh f_camera_MP os displaySize
  <dbl>   <dbl>          <dbl> <fct>      <dbl>      <dbl> <fct>      <dbl>
1 659.    89          3.2 12 GB      5000        16 Andro~      6.7
2 240.    81          2.2 6 GB       5000        16 Andro~      6.59
```

```

3 198.      75                2.4 4 GB          5000          13 Andro~      6.6
4 180.      81                2.2 6 GB          5000          16 Andro~      6.55
5 300.      82                2.6 6 GB          5000          16 Andro~      6.7
6 204.      80                2.2 6 GB          5000           8 Andro~      6.6
# i 1 more variable: displayHz <dbl>

```

General Analysis

After Cleaning:

```

# A tibble: 6 x 9
  price rating `processor GHz)` ram battery_mAh f_camera_MP os displaySize
  <dbl> <dbl>          <dbl> <fct>          <dbl>          <dbl> <fct>          <dbl>
1  659.    89            3.2 12 GB          5000          16 Andro~      6.7
2  240.    81            2.2 6 GB          5000          16 Andro~      6.59
3  198.    75            2.4 4 GB          5000          13 Andro~      6.6
4  180.    81            2.2 6 GB          5000          16 Andro~      6.55
5  300.    82            2.6 6 GB          5000          16 Andro~      6.7
6  204.    80            2.2 6 GB          5000           8 Andro~      6.6
# i 1 more variable: displayHz <dbl>


```

price	rating	processor GHz)	ram
Min. : 88.71	Min. :62.00	Min. :1.600	8 GB :231
1st Qu.: 191.81	1st Qu.:78.00	1st Qu.:2.200	6 GB :144
Median : 275.73	Median :82.00	Median :2.400	4 GB : 81
Mean : 345.50	Mean :81.37	Mean :2.526	12 GB : 38
3rd Qu.: 395.62	3rd Qu.:85.00	3rd Qu.:2.900	16 GB : 7
Max. :2877.34	Max. :89.00	Max. :3.220	3 GB : 3
			(Other): 2

battery_mAh	f_camera_MP	os	displaySize
Min. : 3095	Min. : 1.00	Android v12 :255	Min. :5.900
1st Qu.: 4600	1st Qu.:13.00	Android v11 :150	1st Qu.:6.500
Median : 5000	Median :16.00	Android v13 : 68	Median :6.600
Mean : 4934	Mean :18.16	Android v10 : 17	Mean :6.593
3rd Qu.: 5000	3rd Qu.:20.00	Android v10.0: 4	3rd Qu.:6.670
Max. :22000	Max. :60.00	iOS v15 : 3	Max. :6.950
		(Other) : 9	

displayHz
Min. : 90.0
1st Qu.: 90.0
Median :120.0

Mean :110.5
 3rd Qu.:120.0
 Max. :240.0



Second Dataset

The first dataset looked like this before processing

```
# A tibble: 6 x 12
  Brands Models Colors Memory Storage Camera Rating `Selling Price`
  <chr> <chr> <chr> <chr> <chr> <chr> <dbl> <dbl>
1 SAMSUNG GALAXY M31S Mirage Black 8 GB 128 GB Yes 4.3 19330
2 Nokia 3.2 Steel 2 GB 16 GB Yes 3.8 10199
3 realme C2 Diamond Black 2 GB <NA> Yes 4.4 6999
4 Infinix Note 5 Ice Blue 4 GB 64 GB Yes 4.2 12999
5 Apple iPhone 11 Black 4GB 64 GB Yes 4.6 49900
6 GIONEE L800 Black 8 MB 16 MB Yes 4 2199
# i 4 more variables: `Original Price` <dbl>, Mobile <chr>, Discount <dbl>,
# `discount percentage` <dbl>
```


It is also a tabular data set with information on mobile phones. This data set differs from the first because it has less columns that are useful for predicting price but it has more rows.

Cleaning P1

To start we removed unneeded columns. These were models, Camera, selling price, mobile, discount, and discount percentage. We then make all the column names lowercase. We then made all the data in the colors and brands columns lowercase. We then removed the underscore from the original_price column name. We then converted the price to dollars. We then made the memory, brands, and storage columns factors.

```
# A tibble: 6 x 5
  brands memory storage rating original_price
  <fct>   <fct>   <fct>   <dbl>         <dbl>
1 samsung 8 GB   128 GB   4.3           252.
2 nokia   2 GB   16 GB    3.8           122.
3 infinix 4 GB   64 GB    4.2           156.
4 apple   4GB    64 GB    4.6           599.
5 gionee  8 MB   16 MB    4             26.4
6 apple   3 GB   64 GB    4.6           575.
```

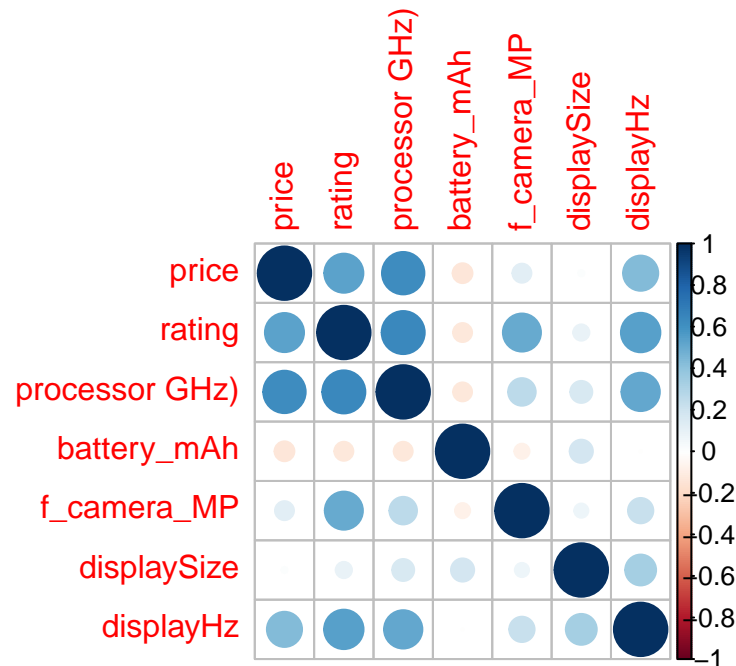
General Analysis

After cleaning:

```
# A tibble: 6 x 5
  brands memory storage rating original_price
  <fct>   <fct>   <fct>   <dbl>         <dbl>
1 samsung 8 GB   128 GB   4.3           252.
2 nokia   2 GB   16 GB    3.8           122.
3 infinix 4 GB   64 GB    4.2           156.
4 apple   4GB    64 GB    4.6           599.
5 gionee  8 MB   16 MB    4             26.4
6 apple   3 GB   64 GB    4.6           575.
```

brands	memory	storage	rating	original_price
samsung:685	4 GB :711	64 GB :757	Min. :2.300	Min. : 12.0
apple :319	3 GB :479	128 GB :720	1st Qu.:4.100	1st Qu.: 124.7
realme :281	6 GB :444	32 GB :545	Median :4.300	Median : 195.6
oppo :251	2 GB :376	16 GB :312	Mean :4.241	Mean : 319.9

xiaomi	:191	8 GB	:326	256 GB	:216	3rd Qu.:4.400	3rd Qu.: 360.0
nokia	:184	1 GB	:193	8 GB	:133	Max. :5.000	Max. :2280.0
(Other)	:986	(Other):368	(Other):214				



Model Creation: Predicting Price

To start with our introductory model, we have decided on a multi-linear regression model. This will set a foundation for the model complex models all predicting price. In the future, we plan to extend upon this will stepwise regression, and a neural network.

Split Test and Train Data:

===== ### Model Creation: Predicting Price

To start with our introductory model, we have decided on a multi-linear regression model. This will set a foundation for the more complex models all predicting price. In the future, we plan to extend upon this will stepwise regression and a neural network.

Fit Two Linear Regression Models:

Summary of Models:

Call:

```
lm(formula = price ~ ., data = df)
```

Residuals:

Min	1Q	Median	3Q	Max
-228.09	-55.93	-14.42	40.95	665.74

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-5.080e+02	3.026e+02	-1.679	0.093801	.
rating	1.754e+01	1.923e+00	9.121	< 2e-16	***
`processor GHz)`	1.423e+02	1.607e+01	8.855	< 2e-16	***
ram16 GB	4.169e+01	4.240e+01	0.983	0.326042	
ram18 GB	3.340e+02	7.465e+01	4.474	9.58e-06	***
ram3 GB	7.775e+01	7.126e+01	1.091	0.275799	
ram4 GB	-5.905e+01	2.889e+01	-2.044	0.041487	*
ram6 GB	-1.195e+02	2.246e+01	-5.320	1.59e-07	***
ram8 GB	-9.567e+01	1.894e+01	-5.052	6.22e-07	***
battery_mAh	-4.114e-04	5.388e-03	-0.076	0.939174	
f_camera_MP	-2.566e+00	5.697e-01	-4.505	8.34e-06	***
osAndroid v10.0	2.238e+01	5.705e+01	0.392	0.695058	
osAndroid v11	-8.211e+01	2.645e+01	-3.104	0.002020	**
osAndroid v12	-6.965e+01	2.570e+01	-2.710	0.006966	**
osAndroid v13	-5.858e+01	2.830e+01	-2.070	0.039017	*
osEMUI v12	-8.082e+01	1.050e+02	-0.770	0.441879	
osHarmony v2.0	-1.405e+02	1.058e+02	-1.328	0.184904	
osHarmonyOS	8.734e+01	1.074e+02	0.813	0.416422	
osHarmonyOS v2.0	-1.582e+01	1.046e+02	-0.151	0.879849	
osHongmeng OS v3.0	2.303e+03	1.064e+02	21.642	< 2e-16	***
osHongmeng OS v4.0	6.622e+02	1.049e+02	6.311	6.29e-10	***
osiOS v15	1.266e+03	6.576e+01	19.258	< 2e-16	***
osiOS v15.0	9.982e+02	6.767e+01	14.750	< 2e-16	***
displaySize	-1.389e+02	3.901e+01	-3.562	0.000405	***
displayHz	1.487e+00	3.293e-01	4.517	7.89e-06	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 101.1 on 481 degrees of freedom

Multiple R-squared: 0.8629, Adjusted R-squared: 0.856

F-statistic: 126.1 on 24 and 481 DF, p-value: < 2.2e-16

Call:

```
lm(formula = original_price ~ ., data = df2)
```

Residuals:

Min	1Q	Median	3Q	Max
-639.24	-77.98	-8.29	46.11	1462.30

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1488.684	109.378	13.610	< 2e-16	***
brandsasus	-491.833	24.253	-20.279	< 2e-16	***
brandsgionee	-517.992	25.129	-20.613	< 2e-16	***
brandsgoogle pixel	-2.467	35.188	-0.070	0.944121	
brandshtc	-281.096	29.623	-9.489	< 2e-16	***
brandsinfinix	-582.695	21.739	-26.804	< 2e-16	***
brandsiqoo	-604.930	77.309	-7.825	7.12e-15	***
brandslenovo	-504.803	24.749	-20.397	< 2e-16	***
brandslg	-398.468	25.380	-15.700	< 2e-16	***
brandsmotorola	-468.750	24.288	-19.300	< 2e-16	***
brandsnokia	-482.104	23.818	-20.241	< 2e-16	***
brandsoppo	-552.058	20.076	-27.498	< 2e-16	***
brandspoco	-606.829	26.194	-23.167	< 2e-16	***
brandsrealme	-612.254	19.535	-31.341	< 2e-16	***
brandssamsung	-426.849	18.813	-22.689	< 2e-16	***
brandsvivo	-566.371	23.302	-24.305	< 2e-16	***
brandsxiaomi	-563.306	20.867	-26.995	< 2e-16	***
memory1.5 GB	-7.670	34.218	-0.224	0.822665	
memory10 MB	-15.231	141.636	-0.108	0.914373	
memory100 MB	-61.859	288.988	-0.214	0.830520	
memory12 GB	469.002	37.802	12.407	< 2e-16	***
memory128 MB	-18.251	196.796	-0.093	0.926115	
memory153 MB	-1309.630	183.502	-7.137	1.21e-12	***
memory16 GB	731.421	172.211	4.247	2.23e-05	***
memory16 MB	-9.389	113.022	-0.083	0.933801	
memory2 GB	-4.512	20.620	-0.219	0.826795	
memory2 MB	-1.560	157.219	-0.010	0.992083	
memory3 GB	64.776	22.653	2.859	0.004275	**
memory30 MB	49.614	213.551	0.232	0.816301	
memory32 MB	3.628	116.866	0.031	0.975234	
memory4 GB	143.757	25.813	5.569	2.80e-08	***
memory4 MB	5.756	120.911	0.048	0.962037	
memory46 MB	-11.058	167.606	-0.066	0.947403	
memory4GB	259.285	27.958	9.274	< 2e-16	***

memory512 MB	-11.839	46.844	-0.253	0.800492	
memory6 GB	203.957	27.609	7.387	1.96e-13	***
memory64 MB	8.718	165.345	0.053	0.957955	
memory768 MB	-22.954	80.677	-0.285	0.776034	
memory8 GB	271.173	29.204	9.285	< 2e-16	***
memory8 MB	13.266	126.205	0.105	0.916295	
storage10 MB	-1211.900	214.648	-5.646	1.81e-08	***
storage100 MB	-1301.445	246.001	-5.290	1.31e-07	***
storage128 GB	-1133.096	70.270	-16.125	< 2e-16	***
storage128 MB	-1238.420	186.657	-6.635	3.88e-11	***
storage129 GB	-1180.067	120.111	-9.825	< 2e-16	***
storage130 GB	-1168.477	138.156	-8.458	< 2e-16	***
storage140 MB	-1172.078	189.277	-6.192	6.78e-10	***
storage153 MB	NA	NA	NA	NA	
storage16 GB	-1160.439	73.034	-15.889	< 2e-16	***
storage16 MB	-1227.396	146.969	-8.351	< 2e-16	***
storage2 MB	-1299.064	125.149	-10.380	< 2e-16	***
storage256 GB	-892.080	70.121	-12.722	< 2e-16	***
storage256 MB	-1144.196	155.237	-7.371	2.22e-13	***
storage32 GB	-1158.027	71.743	-16.141	< 2e-16	***
storage4 GB	-1196.870	85.441	-14.008	< 2e-16	***
storage4 MB	-1260.412	144.187	-8.742	< 2e-16	***
storage512 GB	-392.645	73.653	-5.331	1.05e-07	***
storage512 MB	-1224.462	147.505	-8.301	< 2e-16	***
storage64 GB	-1176.786	70.468	-16.700	< 2e-16	***
storage64 MB	-1225.131	170.344	-7.192	8.13e-13	***
storage8 GB	-1186.759	76.455	-15.522	< 2e-16	***
storage8 MB	-1265.865	183.724	-6.890	6.85e-12	***
storageExpandable Upto 16 GB	-1285.055	149.474	-8.597	< 2e-16	***
storageExpandable Upto 32 GB	-1208.655	136.399	-8.861	< 2e-16	***
rating	64.971	17.707	3.669	0.000248	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 167.4 on 2833 degrees of freedom

Multiple R-squared: 0.7616, Adjusted R-squared: 0.7563

F-statistic: 143.6 on 63 and 2833 DF, p-value: < 2.2e-16

Initial Thoughts:

We would expect to see price increase as the newer specs for phones are released and are put onto the market. Thus, we would expect things like memory, storage, ram, and other specs to be significant predictors for both of our data sets. As we can see this is the true.

One noteworthy observation we noticed in both models in the high level of Adjusted R-Squared(.85 and .75). This can indicate over fitting, or it can indicate an accurate model. This will need to be verified on the test data with some prediction tests after our interpretations.

Interpreting our coefficients:

For the first model,

Intercept: When all other predictor variables are zero, the estimated price of the product is approximately $-\$5.080e+02$.

Rating: For every one-unit increase in the rating, the price is estimated to increase by approximately $\$1.754e+01$, holding all other variables constant.

Processor(GHz): For every one-unit increase in the processor GHz, the price is estimated to increase by approximately $\$1.423e+02$, holding all other variables constant.

RAM(16GB): Phones with 16GB of RAM are estimated to have a price increase of approximately $\$4.169e+01$ compared to the reference category, holding all other variables constant.

RAM(18GB): Phones with 18GB of RAM are estimated to have a price increase of approximately $\$3.340e+02$ compared to the reference category, holding all other variables constant.

RAM(3GB): Phones with 3GB of RAM are estimated to have a price increase of approximately $\$7.775e+01$ compared to the reference category, holding all other variables constant.

RAM(4GB): Phones with 4GB of RAM are estimated to have a price decrease of approximately $\$5.905e+01$ compared to the reference category, holding all other variables constant.

RAM(6GB): Phones with 6GB of RAM are estimated to have a price decrease of approximately $\$1.195e+02$ compared to the reference category, holding all other variables constant.

Ram(8GB): Phones with 6GB of RAM are estimated to have a price decrease of approximately $\$9.567e+01$ compared to the reference category, holding all other variables constant.

Battery(mAH): For every one-unit increase in battery mAh, the price is estimated to decrease by approximately $\$4.114e-04$, holding all other variables constant.

Camera(MP): For every one-unit increase in front camera megapixels, the price is estimated to decrease by approximately $\$2.566$, holding all other variables constant.

Operating System: Devices with Android v10.0: On average, devices with Android v10.0 have prices that are $\$22.38$ higher than the reference category, holding all other variables constant. Android v11: On average, devices with Android v11 have prices that are $\$82.11$ lower than the reference category, holding all other variables constant. Android v12: On average, devices with Android v12 have prices that are $\$69.65$ lower than the reference category, holding all other variables constant. Android v13: On average, devices with Android v13 have prices that are $\$58.58$ lower than the reference category, holding all other variables constant. EMUI v12: On average, devices with EMUI v12 have prices that are $\$80.82$ lower than the reference category, holding all other variables constant. Harmony v2.0: On average, devices with Harmony v2.0

have prices that are \$140.50 lower than the reference category, holding all other variables constant. HarmonyOS: On average, devices with HarmonyOS have prices that are \$87.34 higher than the reference category, holding all other variables constant. HarmonyOS v2.0: On average, devices with HarmonyOS v2.0 have prices that are \$15.82 lower than the reference category, holding all other variables constant. Hongmeng OS v3.0: On average, devices with Hongmeng OS v3.0 have prices that are \$2303 higher than the reference category, holding all other variables constant. Hongmeng OS v4.0: On average, devices with Hongmeng OS v4.0 have prices that are \$662.20 higher than the reference category, holding all other variables constant. iOS v15 and iOS v15.0: On average, devices with iOS v15 have prices that are \$1266 higher than the reference category, and devices with iOS v15.0 have prices that are \$998.20 higher than the reference category, holding all other variables constant.

Display size: For every one-unit increase in display size, the price is estimated to decrease by approximately $1.389e+02$, holding all other variables constant.

Display Hz: For every one-unit increase in display size, the price is estimated to decrease by approximately \$1.487, holding all other variables constant.

For the second model,

Operating System: Phones of the brands ASUS, Gionee, Google Pixel, HTC, Infinix, IQOO, Lenovo, LG, Motorola, Nokia, OPPO, POCO, Realme, Samsung, Vivo, and Xiaomi have an estimated price change of -491.83, -517.99, -2.47, -281.1, -582.7, -604.93, -504.80, -398.47, -468.75, -482.1, -552.06, -606.83, -612.25, -426.85, -566.37, and -563.31 respectively in dollars compared to the reference category holding all other variables constant.

Memory: Phones with RAM memory capacity of 1.5GB, 10MB, 100 MB, 12GB, 128MB, 153MB, 16GB, 16 MB, 2GB, 2MB, 3GB, 30MB, 32MB, 4GB, 46MB, 4G, 512MB, 6GB, 64MB, 768MB, 8GB, and 8MB have an estimated price change of -7.67, -15.23, -61.86, 469.00, -18.25, -1309.63, 731.42, -9.39, -4.51, -1.56, 64.78, 49.61, 3.63, 143.76, 5.76, -11.06, 259.28, -11.84, 203.96, 8.72, -22.95, 271.17, and 13.27 respectively in dollars compared to the reference category holding all other variables constant.

Storage: Phones with storage capacity of 10MB, 100MB, 128GB, 128MB, 129GB, 130GB, 140MB, 153MB, 16GB, 16MB, 2MB, 256GB, 256MB, 32GB, 4GB, 4Mb, 512GB, 512MB, 64GB, 64MB, 8 GB, and 8 MB have an estimated price change of -1211.90, -1301.45, 1133.10, -1238.42, -1180.07, -1168.48, -1172.08, NULL(due to multicollinearity), 1160.44, -1227.40, -1299.06, 892.08, -1144.20, 1158.03, 1196.87, -1260.41, 392.65, -1224.56, 1176.79, -1225.13, 1186.76, and -1265.87 respectively in dollars when compared to the reference category holding all other predictors constant. Rating: For every one unit increase in rating we see a price change of \$64.97 when holding all other predictors constant.

Rating: For every one-unit increase in rating, the price is estimated to increase by approximately \$119.62, holding all other variables constant.

=====

```
[1] 98.61083
```

```
[1] 165.5091
```


```
[1] 345.4993
```

```
[1] 319.8801
```


As you see, we have a RMSE of 98.6 and 165.5 for both data sets respectively. On average, this means our model is off by about \$98.6 and \$165.5. Since the average price of the phones is around \$345.50 and \$319.88, our model is not too inaccurate, but it can be improved.

The RMSE variation is likely caused by the heavy right-skewness in the price variable of more expensive phones. More expensive prices leads to an exponential decay of more expensive components in a thinner market; thus, leading to an exponentially distribution price vector column.


Boxplot of Price Variable(Df1)




Boxplot of Continuous Variable(DF2)



Lasso and Ridge:





Stepwise:

Start: AIC=5653.61
price ~ 1

	Df	Sum of Sq	RSS	AIC
+ os	12	20405396	15476192	5252.1
+ `processor GHz)`	1	14201208	21680381	5400.7
+ rating	1	10260354	25621234	5485.2
+ ram	6	8080244	27801345	5536.5
+ displayHz	1	6707698	29173891	5550.9
+ battery_mAh	1	677355	35204234	5646.0
+ f_camera_MP	1	570140	35311449	5647.5
<none>			35881589	5653.6
+ displaySize	1	5945	35875644	5655.5

Step: AIC=5252.11
price ~ os

	Df	Sum of Sq	RSS	AIC
+ rating	1	7422062	8054130	4923.6
+ `processor GHz)`	1	7269263	8206930	4933.1
+ ram	6	6360285	9115908	4996.3

+ displayHz	1	4136106	11340086	5096.8
+ f_camera_MP	1	746681	14729511	5229.1
+ displaySize	1	63509	15412684	5252.0
<none>			15476192	5252.1
+ battery_mAh	1	59302	15416890	5252.2

Step: AIC=4923.63
price ~ os + rating

	Df	Sum of Sq	RSS	AIC
+ `processor GHz)`	1	1646077	6408053	4809.9
+ ram	6	1586293	6467838	4824.6
+ displayHz	1	561247	7492883	4889.1
+ f_camera_MP	1	338351	7715779	4903.9
<none>			8054130	4923.6
+ battery_mAh	1	5888	8048242	4925.3
+ displaySize	1	4518	8049613	4925.3

Step: AIC=4809.94
price ~ os + rating + `processor GHz)`

	Df	Sum of Sq	RSS	AIC
+ ram	6	979346	5428708	4738.0
+ displayHz	1	240046	6168008	4792.6
+ f_camera_MP	1	225907	6182146	4793.8
<none>			6408053	4809.9
+ displaySize	1	19254	6388799	4810.4
+ battery_mAh	1	16034	6392019	4810.7

Step: AIC=4738.02
price ~ os + rating + `processor GHz)` + ram

	Df	Sum of Sq	RSS	AIC
+ f_camera_MP	1	249163	5179544	4716.2
+ displayHz	1	155134	5273574	4725.4
+ displaySize	1	51875	5376833	4735.2
<none>			5428708	4738.0
+ battery_mAh	1	319	5428389	4740.0

Step: AIC=4716.25
price ~ os + rating + `processor GHz)` + ram + f_camera_MP

	Df	Sum of Sq	RSS	AIC
--	----	-----------	-----	-----

```

+ displayHz      1      125900 5053644 4705.8
+ displaySize    1       50356 5129189 4713.3
<none>                               5179544 4716.2
+ battery_mAh    1        872 5178672 4718.2

```

Step: AIC=4705.8

```

price ~ os + rating + `processor GHz)` + ram + f_camera_MP +
      displayHz

```

```

          Df Sum of Sq      RSS      AIC
+ displaySize 1      133193 4920452 4694.3
<none>                               5053644 4705.8
+ battery_mAh 1        3476 5050168 4707.4

```

Step: AIC=4694.28

```

price ~ os + rating + `processor GHz)` + ram + f_camera_MP +
      displayHz + displaySize

```

```

          Df Sum of Sq      RSS      AIC
<none>                               4920452 4694.3
+ battery_mAh 1        59.627 4920392 4696.3

```

Start: AIC=33757.1

```

original_price ~ 1

```

```

          Df Sum of Sq      RSS      AIC
+ storage 24 178005343 154834932 31588
+ brands  16 152690410 180149866 32011
+ memory  23 136380092 196460183 32276
+ rating   1  67635920 265204355 33101
<none>                               332840275 33757

```

Step: AIC=31588.01

```

original_price ~ storage

```

```

          Df Sum of Sq      RSS      AIC
+ brands 16  66089397  88745535 30008
+ memory 22  27970150 126864782 31055
+ rating  1   7926316 146908616 31438
<none>                               154834932 31588

```

Step: AIC=30007.58

original_price ~ storage + brands

	Df	Sum of Sq	RSS	AIC
+ memory	22	9010140	79735395	29741
+ rating	1	602183	88143352	29990
<none>			88745535	30008

Step: AIC=29741.43

original_price ~ storage + brands + memory

	Df	Sum of Sq	RSS	AIC
+ rating	1	377153	79358242	29730
<none>			79735395	29741

Step: AIC=29729.7

original_price ~ storage + brands + memory + rating

Start: AIC=4696.28

price ~ rating + `processor GHz)` + ram + battery_mAh + f_camera_MP +
os + displaySize + displayHz

	Df	Sum of Sq	RSS	AIC
- battery_mAh	1	60	4920452	4694.3
<none>			4920392	4696.3
- displaySize	1	129776	5050168	4707.4
- f_camera_MP	1	207604	5127996	4715.2
- displayHz	1	208756	5129148	4715.3
- `processor GHz)`	1	802116	5722508	4770.7
- ram	6	946025	5866417	4773.3
- rating	1	850993	5771385	4775.0
- os	12	13374882	18295274	5336.8

Step: AIC=4694.28

price ~ rating + `processor GHz)` + ram + f_camera_MP + os +
displaySize + displayHz

	Df	Sum of Sq	RSS	AIC
<none>			4920452	4694.3
- displaySize	1	133193	5053644	4705.8
- f_camera_MP	1	207549	5128001	4713.2
- displayHz	1	208737	5129189	4713.3
- `processor GHz)`	1	810063	5730514	4769.4

- ram	6	960551	5881003	4772.5
- rating	1	852479	5772930	4773.1
- os	12	13530876	18451328	5339.1

Start: AIC=29729.7

original_price ~ brands + memory + storage + rating

	Df	Sum of Sq	RSS	AIC
<none>			79358242	29730
- rating	1	377153	79735395	29741
- memory	22	8785110	88143352	29990
- storage	23	28630931	107989173	30576
- brands	16	44457006	123815248	30986

Neural Network:

Epoch 1/50

Train metrics: Loss: 112506.5469 - Acc: 0

Epoch 2/50

Train metrics: Loss: 85766.7578 - Acc: 0

Epoch 3/50

Train metrics: Loss: 79610.3672 - Acc: 0

Epoch 4/50

Train metrics: Loss: 81697.3203 - Acc: 0

Epoch 5/50

Train metrics: Loss: 80163.8359 - Acc: 0

Epoch 6/50

Train metrics: Loss: 79525.0781 - Acc: 0

Epoch 7/50

Train metrics: Loss: 65035.8555 - Acc: 0

Epoch 8/50

Train metrics: Loss: 80852.9844 - Acc: 0

Epoch 9/50

Train metrics: Loss: 75625.9766 - Acc: 0

Epoch 10/50

Train metrics: Loss: 78820.7578 - Acc: 0

Epoch 11/50

Train metrics: Loss: 74856.2812 - Acc: 0

Epoch 12/50

Train metrics: Loss: 78154.7969 - Acc: 0

Epoch 13/50

Train metrics: Loss: 72478.1719 - Acc: 0

Epoch 14/50
Train metrics: Loss: 80209.3281 - Acc: 0
Epoch 15/50
Train metrics: Loss: 76235.0625 - Acc: 0
Epoch 16/50
Train metrics: Loss: 78544.0078 - Acc: 0
Epoch 17/50
Train metrics: Loss: 77271.5938 - Acc: 0
Epoch 18/50
Train metrics: Loss: 76881.2109 - Acc: 0
Epoch 19/50
Train metrics: Loss: 71989.1406 - Acc: 0
Epoch 20/50
Train metrics: Loss: 75301.0859 - Acc: 0
Epoch 21/50
Train metrics: Loss: 65565.375 - Acc: 0
Epoch 22/50
Train metrics: Loss: 66891.6641 - Acc: 0
Epoch 23/50
Train metrics: Loss: 69964.75 - Acc: 0
Epoch 24/50
Train metrics: Loss: 68333.7266 - Acc: 0
Epoch 25/50
Train metrics: Loss: 63634.2578 - Acc: 0
Epoch 26/50
Train metrics: Loss: 51410.9688 - Acc: 0
Epoch 27/50
Train metrics: Loss: 61619.1523 - Acc: 0
Epoch 28/50
Train metrics: Loss: 63912.418 - Acc: 0
Epoch 29/50
Train metrics: Loss: 54266.207 - Acc: 0
Epoch 30/50
Train metrics: Loss: 42864.8438 - Acc: 0
Epoch 31/50
Train metrics: Loss: 52724.9883 - Acc: 0
Epoch 32/50
Train metrics: Loss: 53360.2852 - Acc: 0
Epoch 33/50
Train metrics: Loss: 55283.2695 - Acc: 0
Epoch 34/50
Train metrics: Loss: 56338.1562 - Acc: 0
Epoch 35/50

Train metrics: Loss: 51082.9453 - Acc: 0
Epoch 36/50
Train metrics: Loss: 51757.2344 - Acc: 0
Epoch 37/50
Train metrics: Loss: 64937.6875 - Acc: 0
Epoch 38/50
Train metrics: Loss: 65497.8164 - Acc: 0
Epoch 39/50
Train metrics: Loss: 56118.0156 - Acc: 0
Epoch 40/50
Train metrics: Loss: 43047.8828 - Acc: 0
Epoch 41/50
Train metrics: Loss: 54793.7383 - Acc: 0
Epoch 42/50
Train metrics: Loss: 52634.1758 - Acc: 0
Epoch 43/50
Train metrics: Loss: 49707.7617 - Acc: 0
Epoch 44/50
Train metrics: Loss: 50425.582 - Acc: 0
Epoch 45/50
Train metrics: Loss: 57680.4844 - Acc: 0
Epoch 46/50
Train metrics: Loss: 61455.3828 - Acc: 0
Epoch 47/50
Train metrics: Loss: 50908.5195 - Acc: 0
Epoch 48/50
Train metrics: Loss: 52184.6172 - Acc: 0
Epoch 49/50
Train metrics: Loss: 47321.5547 - Acc: 0
Epoch 50/50
Train metrics: Loss: 50804.043 - Acc: 0

Epoch 1/50
Train metrics: Loss: 165096.5938 - Acc: 0
Epoch 2/50
Train metrics: Loss: 69674.6094 - Acc: 0
Epoch 3/50
Train metrics: Loss: 35680.7695 - Acc: 0
Epoch 4/50
Train metrics: Loss: 31565.9258 - Acc: 0
Epoch 5/50
Train metrics: Loss: 30040.6953 - Acc: 0

Epoch 6/50
Train metrics: Loss: 28820.0645 - Acc: 0
Epoch 7/50
Train metrics: Loss: 27476.5332 - Acc: 0
Epoch 8/50
Train metrics: Loss: 26603.627 - Acc: 0
Epoch 9/50
Train metrics: Loss: 25764.3359 - Acc: 0
Epoch 10/50
Train metrics: Loss: 25892.1738 - Acc: 0
Epoch 11/50
Train metrics: Loss: 25045.9062 - Acc: 0
Epoch 12/50
Train metrics: Loss: 24550.5312 - Acc: 0
Epoch 13/50
Train metrics: Loss: 23697.8906 - Acc: 0
Epoch 14/50
Train metrics: Loss: 23902.9414 - Acc: 0
Epoch 15/50
Train metrics: Loss: 23255.5156 - Acc: 0
Epoch 16/50
Train metrics: Loss: 23440.5918 - Acc: 0
Epoch 17/50
Train metrics: Loss: 22738.5293 - Acc: 0
Epoch 18/50
Train metrics: Loss: 23325.0371 - Acc: 0
Epoch 19/50
Train metrics: Loss: 22682.6406 - Acc: 0
Epoch 20/50
Train metrics: Loss: 22552.627 - Acc: 0
Epoch 21/50
Train metrics: Loss: 21867 - Acc: 0
Epoch 22/50
Train metrics: Loss: 22269.502 - Acc: 0
Epoch 23/50
Train metrics: Loss: 22073.7207 - Acc: 0
Epoch 24/50
Train metrics: Loss: 21878.0957 - Acc: 0
Epoch 25/50
Train metrics: Loss: 21743.793 - Acc: 0
Epoch 26/50
Train metrics: Loss: 21527.5625 - Acc: 0
Epoch 27/50

Train metrics: Loss: 21715.8223 - Acc: 0
Epoch 28/50
Train metrics: Loss: 21873.252 - Acc: 0
Epoch 29/50
Train metrics: Loss: 21451.8223 - Acc: 0
Epoch 30/50
Train metrics: Loss: 21179.2285 - Acc: 0
Epoch 31/50
Train metrics: Loss: 21086.6211 - Acc: 0
Epoch 32/50
Train metrics: Loss: 21401.8555 - Acc: 0
Epoch 33/50
Train metrics: Loss: 20948.5293 - Acc: 0
Epoch 34/50
Train metrics: Loss: 21148.4258 - Acc: 0
Epoch 35/50
Train metrics: Loss: 21381.3418 - Acc: 0
Epoch 36/50
Train metrics: Loss: 20817.4121 - Acc: 0
Epoch 37/50
Train metrics: Loss: 21304.4316 - Acc: 0
Epoch 38/50
Train metrics: Loss: 21120.2734 - Acc: 0
Epoch 39/50
Train metrics: Loss: 21004.4629 - Acc: 0
Epoch 40/50
Train metrics: Loss: 20604.1328 - Acc: 0
Epoch 41/50
Train metrics: Loss: 20290.3477 - Acc: 0
Epoch 42/50
Train metrics: Loss: 20742.4863 - Acc: 0
Epoch 43/50
Train metrics: Loss: 20452.8125 - Acc: 0
Epoch 44/50
Train metrics: Loss: 20365.6172 - Acc: 0
Epoch 45/50
Train metrics: Loss: 20576.9746 - Acc: 0
Epoch 46/50
Train metrics: Loss: 20645.0684 - Acc: 0
Epoch 47/50
Train metrics: Loss: 20185.3242 - Acc: 0
Epoch 48/50
Train metrics: Loss: 20478.6211 - Acc: 0

Epoch 49/50

Train metrics: Loss: 20762.3398 - Acc: 0

Epoch 50/50

Train metrics: Loss: 20324.2051 - Acc: 0

Summary of results

	Model	Dataset	RMSE
1	Linear	1	98.61083
2	Lasso	1	173.73677
3	Ridge	1	170.54350
4	Forwards	1	98.61142
5	Backwards	1	98.61142
6	NNetwork	1	301.61023
7	Linear	2	165.50906
8	Lasso	2	286.82097
9	Ridge	2	288.88781
10	Forwards	2	165.50906
11	Backwards	2	165.50906
12	NNetwork	2	452.43591